



West Visayas State University

(Formerly Iloilo Normal School)

COLLEGE OF INFORMATION AND COMMUNICATIONS TECHNOLOGY

Luna St., La Paz, Iloilo City 5000

Iloilo, Philippines

* Trunkline: (063) (033) 320-0870 to 78 loc. 1403 * Telefax No.: (033) 320-0879

* Website: www.wvsu.edu.ph

* Email Address: cict@wvsu.edu.ph



Memowee: Face Recognition System for Dementia Patient Care

Project Documentation

A Final Project

Presented to Mr. John Christopher Mateo

Faculty, College of Information and Communications Technology

West Visayas State University

La Paz, Iloilo City

In Partial Fulfillment

of the Requirements for the Course

CCS 248: Artificial Neural Networks

By

Jethro A. Rendon

Joelmar D. Grecia

Axel John E. Nuqui

December 2025





West Visayas State University

(Formerly Iloilo Normal School)

COLLEGE OF INFORMATION AND COMMUNICATIONS TECHNOLOGY

Luna St., La Paz, Iloilo City 5000

Iloilo, Philippines

* Trunkline: (063) (033) 320-0870 to 78 loc. 1403 * Telefax No.: (033) 320-0879

* Website: www.wvsu.edu.ph

* Email Address: cict@wvsu.edu.ph



Problem

Dementia is a progressive neurological disorder that impairs cognitive function, including memory loss and difficulty recognizing familiar people. Patients with dementia often struggle to identify family members, caregivers, and close friends, creating emotional distress for both the patient and their loved ones.

Memowee addresses this challenge by developing a face recognition system that allows dementia patients (or their caregivers) to pre-register known persons with their facial profiles and automatically identify them when they appear. The system displays contextual information (name, relationship, description) to help patients reconnect with their loved ones.

Dataset Preparation

The project uses a **subset of the VGGFace2** dataset (from Kaggle), which contains a large-scale collection of face images with variations in pose, age, lighting, and expression.[3] The subset was extracted and organized into train/validation folders, then renamed as **MEMOWEE_DATASET** for project convenience.

1. MEMOWEE_DATASET Statistics

- Training Samples: 18,400 face image pairs
- Validation Samples: 2,400 face image pairs
- Total identities: Multiple persons across various demographic groups
- Image format: JPEG, resized to 128×128 pixels

2. Data Preprocessing

1. Image Loading: Images are loaded using OpenCV and converted from BGR to RGB color space.
2. Resizing: All images are resized to 128×128 pixels to match the model input.
3. Normalization: Pixel values are normalized to [0, 1] by dividing by 255.
4. Augmentation (during training):
 - a. Random Horizontal Lines
 - b. Random Rotations ($\pm 10\%$)
 - c. Random Zoom ($\pm 10\%$)





West Visayas State University

(Formerly Iloilo Normal School)

COLLEGE OF INFORMATION AND COMMUNICATIONS TECHNOLOGY

Luna St., La Paz, Iloilo City 5000

Iloilo, Philippines

* Trunkline: (063) (033) 320-0870 to 78 loc. 1403 * Telefax No.: (033) 320-0879

* Website: www.wvsu.edu.ph

* Email Address: cict@wvsu.edu.ph



Data Validation and Quality Checks

1. Bias Assessment:

[VGGFace2](#) is a well-established benchmark dataset with diverse identities, ages, and ethnicities, reducing demographic bias. [3]

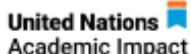
2. Privacy Considerations:

- VGGFace2 is properly licensed for academic research use.
- The system stores 128-dimensional face embeddings (not raw images), reducing privacy exposure.
- Patient profiles are stored locally and access-controlled.

Chosen Deep Neural Network

Siamese Neural Network with Contrastive Learning is chosen because:

- Face verification is inherently a similarity-learning task requiring the network to determine if two faces belong to the same person.[4]
- Siamese networks excel at one-shot and few-shot learning, suitable for scenarios where patients may register only a few photos per person.
- Contrastive loss directly optimizes learned embeddings: same-person pairs are close, different-person pairs are far apart.





West Visayas State University

(Formerly Iloilo Normal School)

COLLEGE OF INFORMATION AND COMMUNICATIONS TECHNOLOGY

Luna St., La Paz, Iloilo City 5000

Iloilo, Philippines

* Trunkline: (063) (033) 320-0870 to 78 loc. 1403 * Telefax No.: (033) 320-0879

* Website: www.wvsu.edu.ph

* Email Address: cict@wvsu.edu.ph



Network Architecture

Embedding Model (Shared between Siamese inputs):

Layer	Output Shape	Parameters	Details
Input	(128, 128, 3)	0	RGB Images
Random Flip	(128, 128, 3)	0	Horizontal flip augmentation
RandomRotation	(128, 128, 3)	0	$\pm 10^\circ$ rotation augmentation
RandomZoom	(128, 128, 3)	0	$\pm 10\%$ zoom augmentation
Conv2D Block 1	(128, 128, 32)	896	32 filters, 3×3 kernel, ReLU, L2 regularizer (1e-4)
MaxPooling2D 1	(64, 64, 32)	0	2×2 pooling
Dropout 1	(64, 64, 32)	0	0.25 dropout
Conv2D Block 2	(64, 64, 64)	18,496	64 filters, 3×3 kernel, ReLU, L2 regularizer (1e-4)
MaxPooling2D 2	(32, 32, 64)	0	2×2 pooling
Dropout 2	(32, 32, 64)	0	0.25 dropout
Conv2D Block 3	(32, 32, 128)	73,856	128 filters, 3×3 kernel, ReLU, L2 regularizer (1e-4)
MaxPooling2D 3	(16, 16, 128)	0	2×2 pooling
Dropout 3	(16, 16, 128)	0	0.25 dropout
Flatten	(32,768)	0	Flattens spatial dimensions
Dense	(128)	4,194,432	128 units, ReLU, L2 regularizer (1e-4)
L2 Normalization	(128)	0	Unit-norm embedding output

Total Parameters: 4,287,680 (16.36 MB)

Trainable Parameters: 4,287,680

Non-Trainable Parameters: 0





West Visayas State University

(Formerly Iloilo Normal School)

COLLEGE OF INFORMATION AND COMMUNICATIONS TECHNOLOGY

Luna St., La Paz, Iloilo City 5000

Iloilo, Philippines

* Trunkline: (063) (033) 320-0870 to 78 loc. 1403 * Telefax No.: (033) 320-0879

* Website: www.wvsu.edu.ph

* Email Address: cict@wvsu.edu.ph



Training Configuration

Loss Function:

Contrastive Loss is used to train the Siamese network

$$L = y \cdot d^2 + (1 - y) \cdot \max(\text{margin} - d, 0)^2$$

Where:

- $y = 1$ for same-person pairs (positive pairs)
- $y = 0$ for different-person pairs (negative pairs)
- d = Euclidean distance between embeddings
- $\text{margin} = 1.0$

This loss encourages small distances for same-person pairs and large distances for different-person pairs.

Optimizer and Learning Rate Strategy

Parameter	Value
Optimizer	Adam
Initial Learning Rate	0.001
Learning Rate Schedule	ReduceLROnPlateau
LR Reduction Factor	0.5 (reduce by 50% when val_loss plateaus)
Patience	3 epochs without improvement
Minimum Learning Rate	1e-6
Loss Monitor	val_loss (validation loss)

Adam is chosen for its adaptive learning rates and efficiency. ReduceLROnPlateau dynamically adjusts the learning rate based on validation performance, improving convergence.

Training Configuration Summary

Parameter	Value
Batch Size	32
Number of Epochs	30
Image Size	128×128 pixels
Data Generator	Custom SiameseDataGenerator (online pair generation)
Training Pairs	18,400
Validation Pairs	2,400
Pair Generation	For each anchor image: 1 positive pair + 1 negative pair per batch





West Visayas State University

(Formerly Iloilo Normal School)

COLLEGE OF INFORMATION AND COMMUNICATIONS TECHNOLOGY

Luna St., La Paz, Iloilo City 5000

Iloilo, Philippines

* Trunkline: (063) (033) 320-0870 to 78 loc. 1403 * Telefax No.: (033) 320-0879

* Website: www.wvsu.edu.ph

* Email Address: cict@wvsu.edu.ph



Hyper Parameter Tuning Records

The model was trained for 30 epochs with continuous monitoring of training and validation metrics.

Key epochs are recorded below:

```

Epoch 1/30
575/575 604s 1s/step - contrastive_accuracy_metric: 0.5440 - loss: 0.2696 - val_contrastive_accuracy_metric: 0.5819 - val_loss: 0.2394 - learning_rate: 0.0010
Epoch 2/30
575/575 602s 1s/step - contrastive_accuracy_metric: 0.5811 - loss: 0.2406 - val_contrastive_accuracy_metric: 0.6110 - val_loss: 0.2282 - learning_rate: 0.0010
Epoch 3/30
575/575 604s 1s/step - contrastive_accuracy_metric: 0.5953 - loss: 0.2327 - val_contrastive_accuracy_metric: 0.6027 - val_loss: 0.2244 - learning_rate: 0.0010
Epoch 4/30
575/575 600s 1s/step - contrastive_accuracy_metric: 0.6022 - loss: 0.2303 - val_contrastive_accuracy_metric: 0.6285 - val_loss: 0.2217 - learning_rate: 0.0010
Epoch 5/30
575/575 571s 993ms/step - contrastive_accuracy_metric: 0.6098 - loss: 0.2268 - val_contrastive_accuracy_metric: 0.6510 - val_loss: 0.2139 - learning_rate: 0.0010
Epoch 6/30
575/575 505s 878ms/step - contrastive_accuracy_metric: 0.6191 - loss: 0.2249 - val_contrastive_accuracy_metric: 0.6438 - val_loss: 0.2108 - learning_rate: 0.0010
Epoch 7/30
575/575 504s 876ms/step - contrastive_accuracy_metric: 0.6224 - loss: 0.2231 - val_contrastive_accuracy_metric: 0.6708 - val_loss: 0.2142 - learning_rate: 0.0010
Epoch 8/30
575/575 501s 871ms/step - contrastive_accuracy_metric: 0.6299 - loss: 0.2196 - val_contrastive_accuracy_metric: 0.6596 - val_loss: 0.2058 - learning_rate: 0.0010
Epoch 9/30
575/575 501s 871ms/step - contrastive_accuracy_metric: 0.6380 - loss: 0.2164 - val_contrastive_accuracy_metric: 0.6706 - val_loss: 0.2070 - learning_rate: 0.0010
Epoch 10/30
575/575 494s 860ms/step - contrastive_accuracy_metric: 0.6401 - loss: 0.2155 - val_contrastive_accuracy_metric: 0.6777 - val_loss: 0.2031 - learning_rate: 0.0010
Epoch 11/30
575/575 493s 857ms/step - contrastive_accuracy_metric: 0.6530 - loss: 0.2127 - val_contrastive_accuracy_metric: 0.6715 - val_loss: 0.2100 - learning_rate: 0.0010
Epoch 12/30
575/575 500s 870ms/step - contrastive_accuracy_metric: 0.6496 - loss: 0.2123 - val_contrastive_accuracy_metric: 0.6719 - val_loss: 0.2003 - learning_rate: 0.0010
Epoch 13/30
575/575 499s 867ms/step - contrastive_accuracy_metric: 0.6588 - loss: 0.2100 - val_contrastive_accuracy_metric: 0.6988 - val_loss: 0.1982 - learning_rate: 0.0010
Epoch 14/30
575/575 506s 880ms/step - contrastive_accuracy_metric: 0.6609 - loss: 0.2084 - val_contrastive_accuracy_metric: 0.6800 - val_loss: 0.1974 - learning_rate: 0.0010
Epoch 15/30
575/575 480s 835ms/step - contrastive_accuracy_metric: 0.6638 - loss: 0.2068 - val_contrastive_accuracy_metric: 0.7079 - val_loss: 0.1959 - learning_rate: 0.0010
Epoch 16/30
575/575 467s 813ms/step - contrastive_accuracy_metric: 0.6691 - loss: 0.2046 - val_contrastive_accuracy_metric: 0.7198 - val_loss: 0.1941 - learning_rate: 0.0010
Epoch 17/30
575/575 500s 870ms/step - contrastive_accuracy_metric: 0.6670 - loss: 0.2074 - val_contrastive_accuracy_metric: 0.6956 - val_loss: 0.1958 - learning_rate: 0.0010
Epoch 18/30
575/575 474s 825ms/step - contrastive_accuracy_metric: 0.6717 - loss: 0.2048 - val_contrastive_accuracy_metric: 0.6975 - val_loss: 0.1971 - learning_rate: 0.0010
Epoch 19/30
575/575 500s 870ms/step - contrastive_accuracy_metric: 0.6723 - loss: 0.2044 - val_contrastive_accuracy_metric: 0.7100 - val_loss: 0.1934 - learning_rate: 0.0010
Epoch 20/30
575/575 509s 886ms/step - contrastive_accuracy_metric: 0.6777 - loss: 0.2029 - val_contrastive_accuracy_metric: 0.7069 - val_loss: 0.1901 - learning_rate: 0.0010
Epoch 21/30
575/575 476s 828ms/step - contrastive_accuracy_metric: 0.6798 - loss: 0.2023 - val_contrastive_accuracy_metric: 0.7144 - val_loss: 0.1887 - learning_rate: 0.0010
Epoch 22/30
575/575 492s 856ms/step - contrastive_accuracy_metric: 0.6787 - loss: 0.2025 - val_contrastive_accuracy_metric: 0.6910 - val_loss: 0.1960 - learning_rate: 0.0010
Epoch 23/30
575/575 487s 847ms/step - contrastive_accuracy_metric: 0.6801 - loss: 0.2013 - val_contrastive_accuracy_metric: 0.7050 - val_loss: 0.1927 - learning_rate: 0.0010
Epoch 24/30
575/575 0s 805ms/step - contrastive_accuracy_metric: 0.6808 - loss: 0.2014
Epoch 24: ReduceLROnPlateau reducing learning rate to 0.000500000237487257.
575/575 479s 833ms/step - contrastive_accuracy_metric: 0.6804 - loss: 0.2017 - val_contrastive_accuracy_metric: 0.6823 - val_loss: 0.1980 - learning_rate: 0.0010
Epoch 25/30
575/575 478s 831ms/step - contrastive_accuracy_metric: 0.6911 - loss: 0.1961 - val_contrastive_accuracy_metric: 0.7319 - val_loss: 0.1831 - learning_rate: 5.0000e-04
Epoch 26/30
575/575 476s 828ms/step - contrastive_accuracy_metric: 0.6948 - loss: 0.1930 - val_contrastive_accuracy_metric: 0.7260 - val_loss: 0.1821 - learning_rate: 5.0000e-04
Epoch 27/30
575/575 478s 831ms/step - contrastive_accuracy_metric: 0.6952 - loss: 0.1922 - val_contrastive_accuracy_metric: 0.7260 - val_loss: 0.1798 - learning_rate: 5.0000e-04
Epoch 28/30
575/575 479s 833ms/step - contrastive_accuracy_metric: 0.6961 - loss: 0.1910 - val_contrastive_accuracy_metric: 0.7219 - val_loss: 0.1789 - learning_rate: 5.0000e-04
Epoch 29/30
575/575 481s 836ms/step - contrastive_accuracy_metric: 0.6980 - loss: 0.1898 - val_contrastive_accuracy_metric: 0.7229 - val_loss: 0.1823 - learning_rate: 5.0000e-04
Epoch 30/30
575/575 498s 866ms/step - contrastive_accuracy_metric: 0.6997 - loss: 0.1897 - val_contrastive_accuracy_metric: 0.7306 - val_loss: 0.1806 - learning_rate: 5.0000e-04

```



West Visayas State University

(Formerly Iloilo Normal School)

COLLEGE OF INFORMATION AND COMMUNICATIONS TECHNOLOGY

Luna St., La Paz, Iloilo City 5000

Iloilo, Philippines

* Trunkline: (063) (033) 320-0870 to 78 loc. 1403 * Telefax No.: (033) 320-0879

* Website: www.wvsu.edu.ph * Email Address: cict@wvsu.edu.ph



Results and Evaluation

Metric	Value
Final Training Accuracy	69.97%
Final Validation Accuracy	73.06%
Final Training Loss	0.1897
Final Validation Loss	0.1806

Interpretation

- Above Requirement:** The project requirement is $\geq 50\text{--}60\%$ accuracy. This model achieves 73.06% validation accuracy, well exceeding the baseline.
- Good generalization:** Validation accuracy slightly exceeds training accuracy, indicating the model generalizes well without overfitting.
- Steady convergence:** Both losses decrease monotonically, confirming effective learning throughout training.
- Learning rate adjustment:** The ReduceLROnPlateau callback reduced the learning rate at epoch 20, contributing to final convergence.



Fig 1. Training and Validation Loss



West Visayas State University

(Formerly Iloilo Normal School)

COLLEGE OF INFORMATION AND COMMUNICATIONS TECHNOLOGY

Luna St., La Paz, Iloilo City 5000

Iloilo, Philippines

* Trunkline: (063) (033) 320-0870 to 78 loc. 1403 * Telefax No.: (033) 320-0879

* Website: www.wvsu.edu.ph

* Email Address: cict@wvsu.edu.ph



6. Application Overview

Memowee's Process Flow

- Registration Phase
 1. A caregiver enters the known person's name, relationship (e.g., "daughter"), and a personal note.
 2. A face image is captured or uploaded.
 3. The backend extracts a 128-dimensional embedding using the trained CNN.
 4. The embedding and profile are stored in a secure database.
- Identificaton Phase
 1. When the person appears before the patient, an image is captured.
 2. The system extracts the embedding and compares it to all registered profiles.
 3. If a match is found (distance < threshold), the person's name, relationship, and note are displayed.
 4. If no match is found, the system returns "unknown."





Screenshots

1. Axel

The screenshots illustrate the process of identifying and registering a person named Axel using the MEMOWEE software.

- Screenshot 1: Identification Phase**
The software displays a video feed of a person. A bounding box highlights the face, and the text "Face detected" appears below it. In the bottom left corner of the video frame, a message box says "Not Recognized". The activity log shows the system is processing the face and analyzing it for recognition.
- Screenshot 2: Registration Phase**
The software displays a video feed of the same person. A bounding box highlights the face, and the text "Face detected" appears below it. A registration dialog box is open, showing two options: "Axel" and "Baby". The "Axel" option is selected. The activity log shows the system is processing the face and preparing it for registration.
- Screenshot 3: Verification Phase**
The software displays a video feed of the person. A bounding box highlights the face, and the text "Axel" appears below it. The activity log shows the system is processing the face and verifying it against the database.



2. Joelmar

The figure consists of three vertically stacked screenshots of a computer application window titled "MEMOWEE".

- Screenshot 1:** Shows a video feed of a person's face. A black rectangular crop box is drawn around the face. A small white box in the bottom-left corner of the video area says "Not Recognized" with a question mark icon and a link "Click here to my log". Below the video, the text "Face detected" is displayed.
- Screenshot 2:** Shows the same video feed. A registration dialog box is overlaid on the video. It contains:
 - A dropdown menu showing "Joelmar" selected, with other options "Father" and "Mother".
 - A text input field containing "Joelmar and really cares about the world!"
 - Two buttons: "Cancel" and "Register".
- Screenshot 3:** Shows the video feed again. The black crop box is still present. A white box in the bottom-left corner now displays "Joelmar" with a small profile picture, "Father", and a link "Click here to my log". Below the video, the text "Joelmar" is displayed.

Each screenshot includes an "Activity Log" panel on the right side showing processing logs and a "Recognition Threshold" slider at the bottom.



3. Jethro

The screenshots illustrate the MEMOWEE software's face recognition workflow:

- Screenshot 1: Identification Phase**
The software displays a video feed of a person's face. A bounding box highlights the detected face, and the text "Face detected" is visible. An overlay message says "Not Recognized". The activity log shows the process: "Processing 1 facetc(s) with threshold 0.5...", "Analyzing face...", "Sending to backend for recognition...", "Face 1 not recognized (distance: NaN)", and "Face identification completed".
- Screenshot 2: Registration Phase**
The software displays a video feed of a person's face. A bounding box highlights the detected face, and the text "Face detected" is visible. An overlay message says "Not Recognized". A registration dialog box is open, prompting the user to "Click 'Identify' to try again". The dialog lists names: "Jethro", "Asawa ni Joeross", and "Ang galing and really cool". The "Register" button is highlighted. The activity log shows: "Analyzing face...", "Sending to backend for recognition...", "Face 1 not recognized (distance: NaN)", "Face identification completed", and "Preparing face for registration...".
- Screenshot 3: Identification Phase (Post-Registration)**
The software displays a video feed of a person's face. A bounding box highlights the detected face, and the text "Jethro" is displayed below it. The activity log shows: "Processing 1 facetc(s) with threshold 0.5...", "Analyzing face...", "Sending to backend for recognition...", "Face 1 identified as: Jethro (distance: 0.20224320848983/0)", and "Face identification completed".